

# Annif at SemEval-2025 Task 5: Traditional XMTC augmented by LLMs

Osma Suominen, Juho Inkinen and Mona Lehtinen / National Library of Finland, University of Helsinki

## LLMs4Subjects task overview

### Given:

- GND subject vocabulary (German monolingual terms)
- bibliographic records (titles and abstracts in German and/or English + GND subject labels) from the TIBKAT database

Predict GND subjects for *unseen* titles and abstracts in the test set - an Extreme Multilabel Text Classification (XMTC) task!

### Variations:

- **all**: full GND (200k subjects), large train set (82k records)
- **tib-core**: GND subset (79k subjects), small train set (42k records)

## Our approach

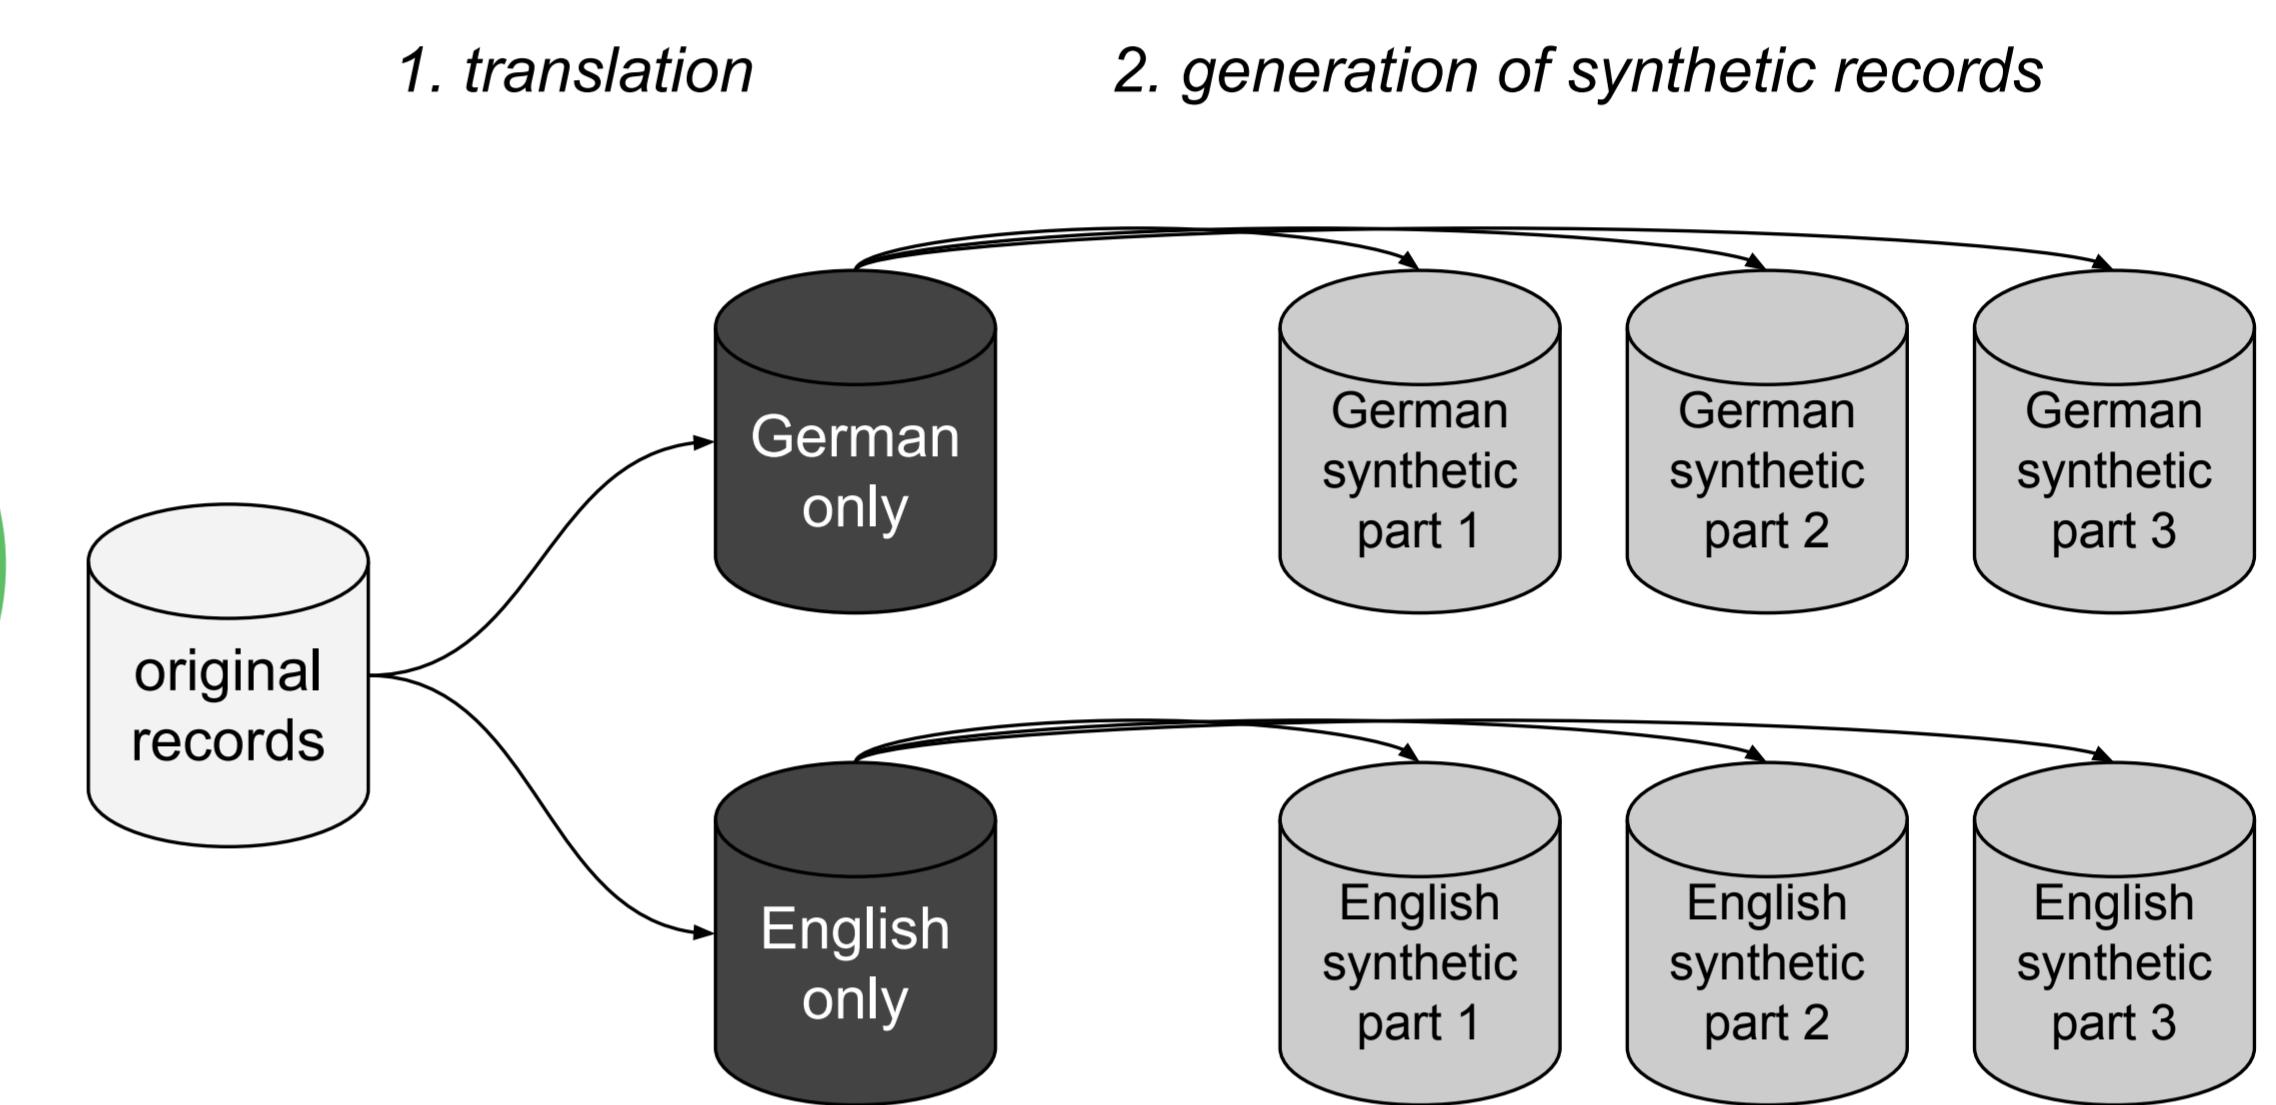
1. Use the **Annif** automated subject indexing toolkit (by us, open source)
2. Pre-process and enhance data set using LLMs
  - translate **GND terms** to English with **GPT-4o-mini**
  - translate **TIBKAT records** to monolingual German/English with **Llama-3.1-8B-Instruct**
  - generate additional **synthetic training data** with **Llama-3.1-8B-Instruct**
3. Train ensemble of XMTC algorithms
  - Omikuji Bonsai (partitioned label trees)
  - MLLM (lexical approach)
  - XTransformer (BERT based XMTC and ranking)
4. Use monolingual (de/en) prediction pipelines in parallel
  - merge their results for final system output

## Quantitative results

Vocab	System	Run#	Ensemble	Lang	development set		test set		
					F1@5	nDCG@10	F1@5	Avg recall	Rank
all	Annif (ours)	1	BM simple	de	0.3174	0.5459	0.3108	0.5736	1 <sup>st</sup>
		2		en	0.3312	0.5677	0.3184	0.5890	
		3		de+en	-	-	0.3376	0.6201	
		4		de	0.3337*	0.5726*	0.3029	0.5447	
		5	BM neural	en	<b>0.3504*</b>	<b>0.6008*</b>	0.3116	0.5599	
		6		de+en	-	-	0.3318	0.6005	
		7		de	0.3263	0.5614	0.3185	0.5859	
		8		en	0.3411	0.5842	0.3276	0.6038	
		9		de+en	-	-	<b>0.3432</b>	<b>0.6295</b>	1 <sup>st</sup>
DUTIR831		3					0.3346	0.6045	2 <sup>nd</sup>
		1					0.3015	0.5856	3 <sup>rd</sup>
		1					0.3231	0.5631	4 <sup>th</sup>
		1					0.2618	0.5302	5 <sup>th</sup>
tib-core	Annif (ours)	1	BM simple	de	0.2821	0.5557	0.2796	0.5285	2 <sup>nd</sup>
		2		en	0.3009	0.5936	0.2984	0.5617	
		3		de+en	-	-	0.3113	0.5824	
		4		de	0.3209*	0.6171*	0.2660	0.4864	
		5	BM neural	en	<b>0.3467*</b>	<b>0.6661*</b>	0.2886	0.5217	
		6		de+en	-	-	0.3043	0.5559	
		7		de	0.2891	0.5684	0.2864	0.5385	
		8		en	0.3079	0.6051	0.3030	0.5719	
		9		de+en	-	-	<b>0.3136</b>	<b>0.5899</b>	2 <sup>nd</sup>
RUC Team		1					<b>0.3271</b>	<b>0.6568</b>	1 <sup>st</sup>
		2					0.2717	0.5794	3 <sup>rd</sup>
		2					0.3153	0.5599	4 <sup>th</sup>
		1					0.2370	0.4976	5 <sup>th</sup>

Table 1: Quantitative evaluation results for the ensemble projects measured against the development and test sets. Top 5 systems included for comparison. Note that Lang refers to the project language, not to the indicated language of the records. \*Unreliable score because the neural ensemble was trained on the development set it was evaluated on.

## Pre-processing steps

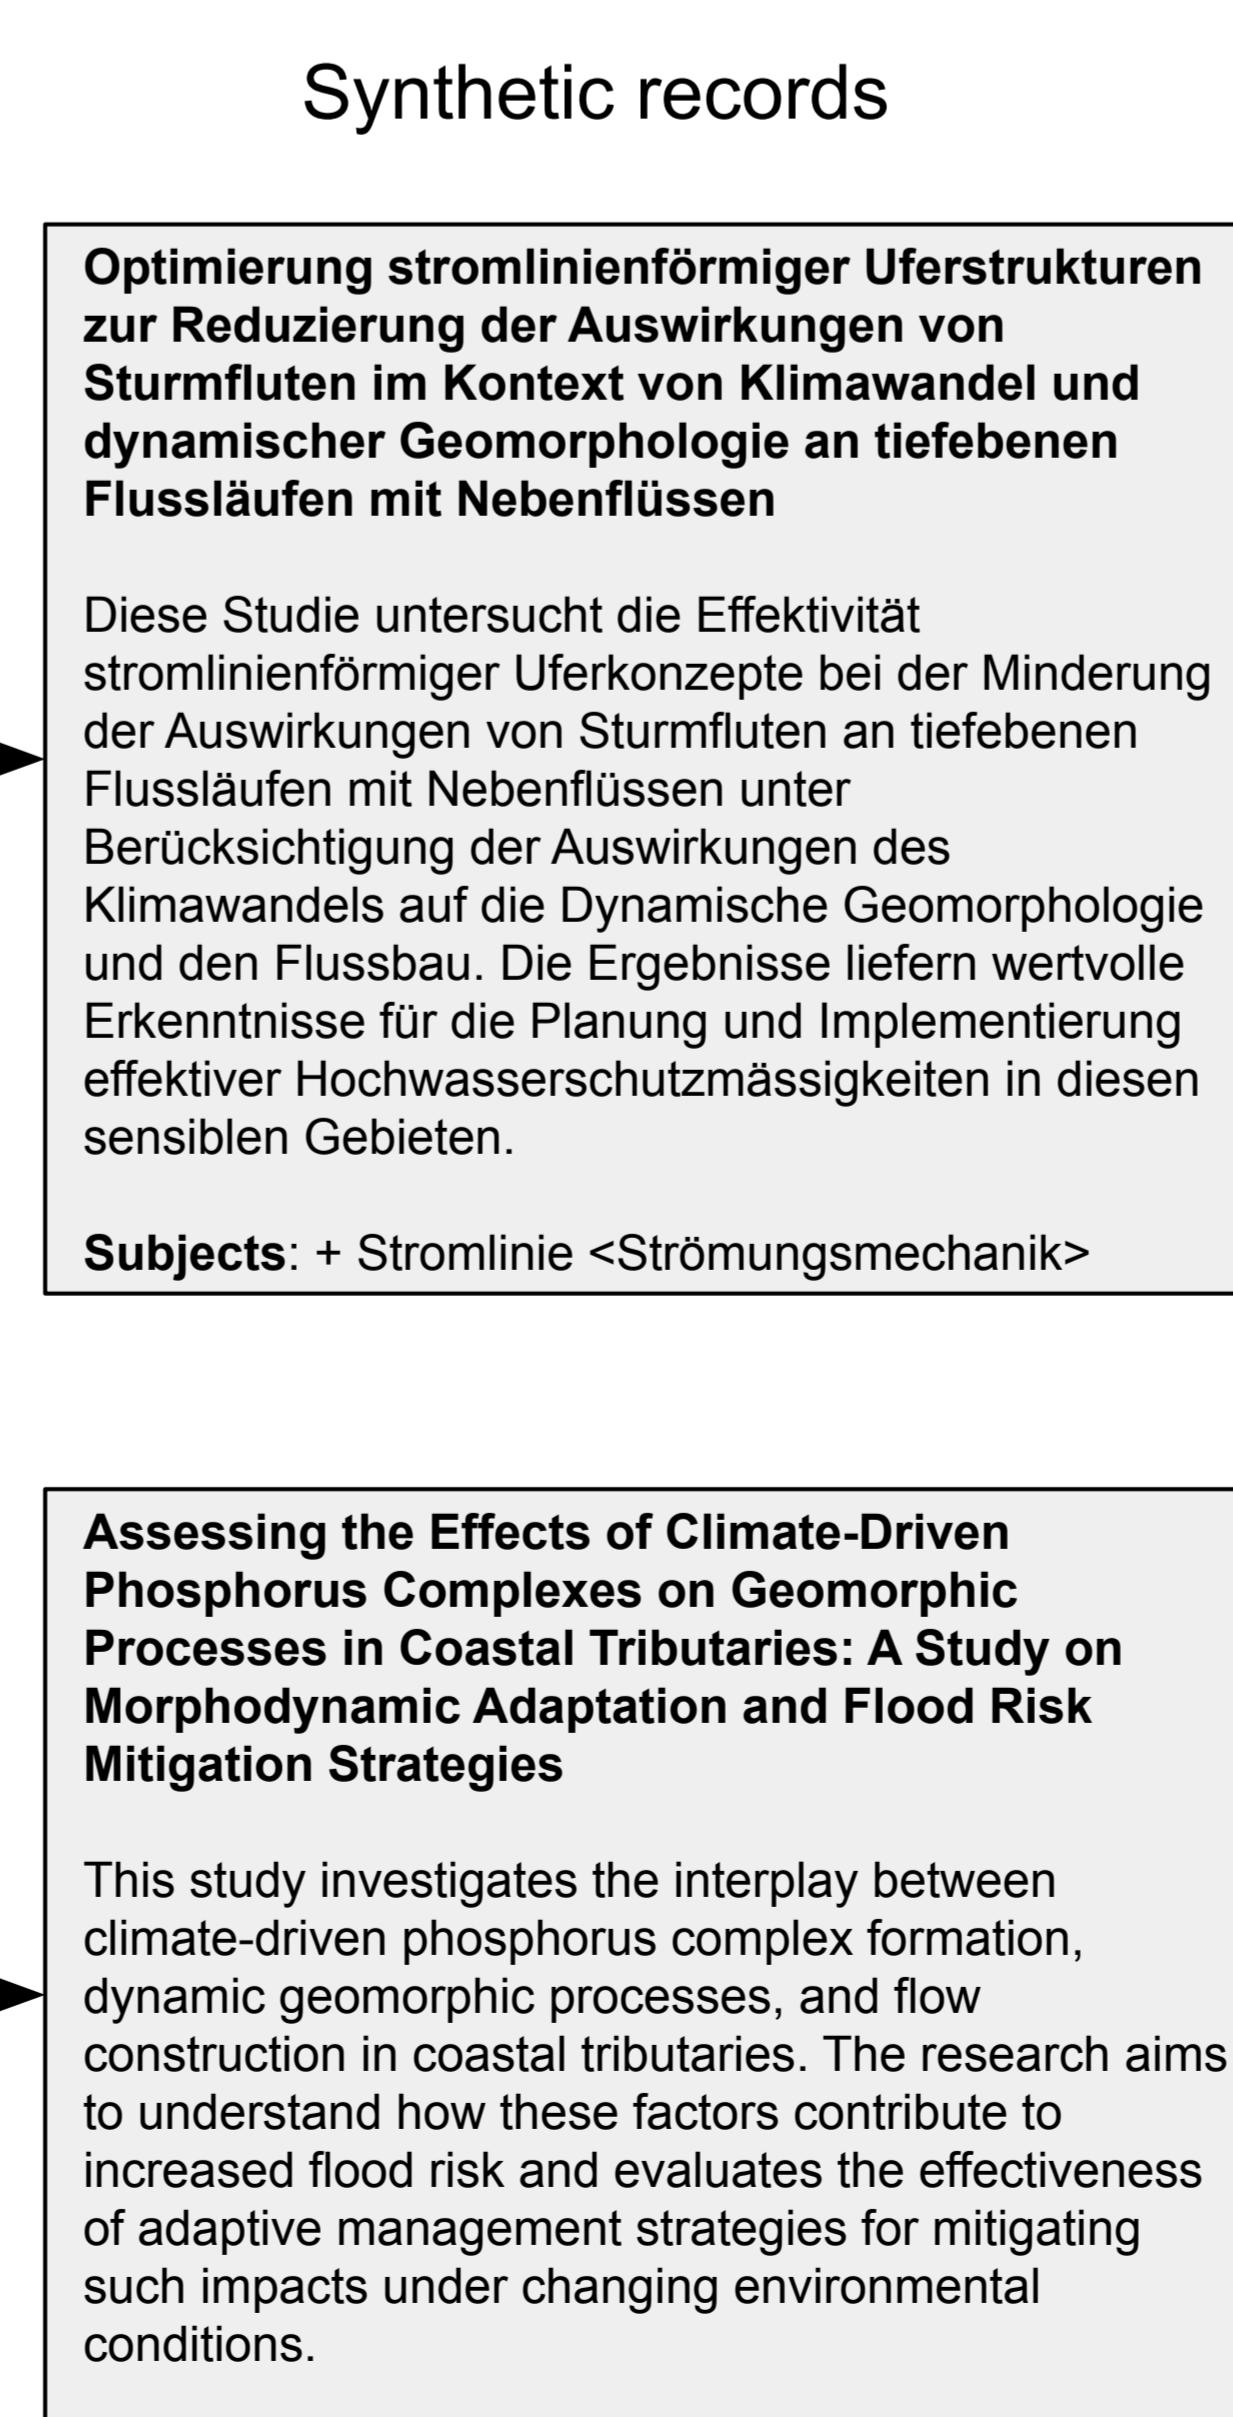
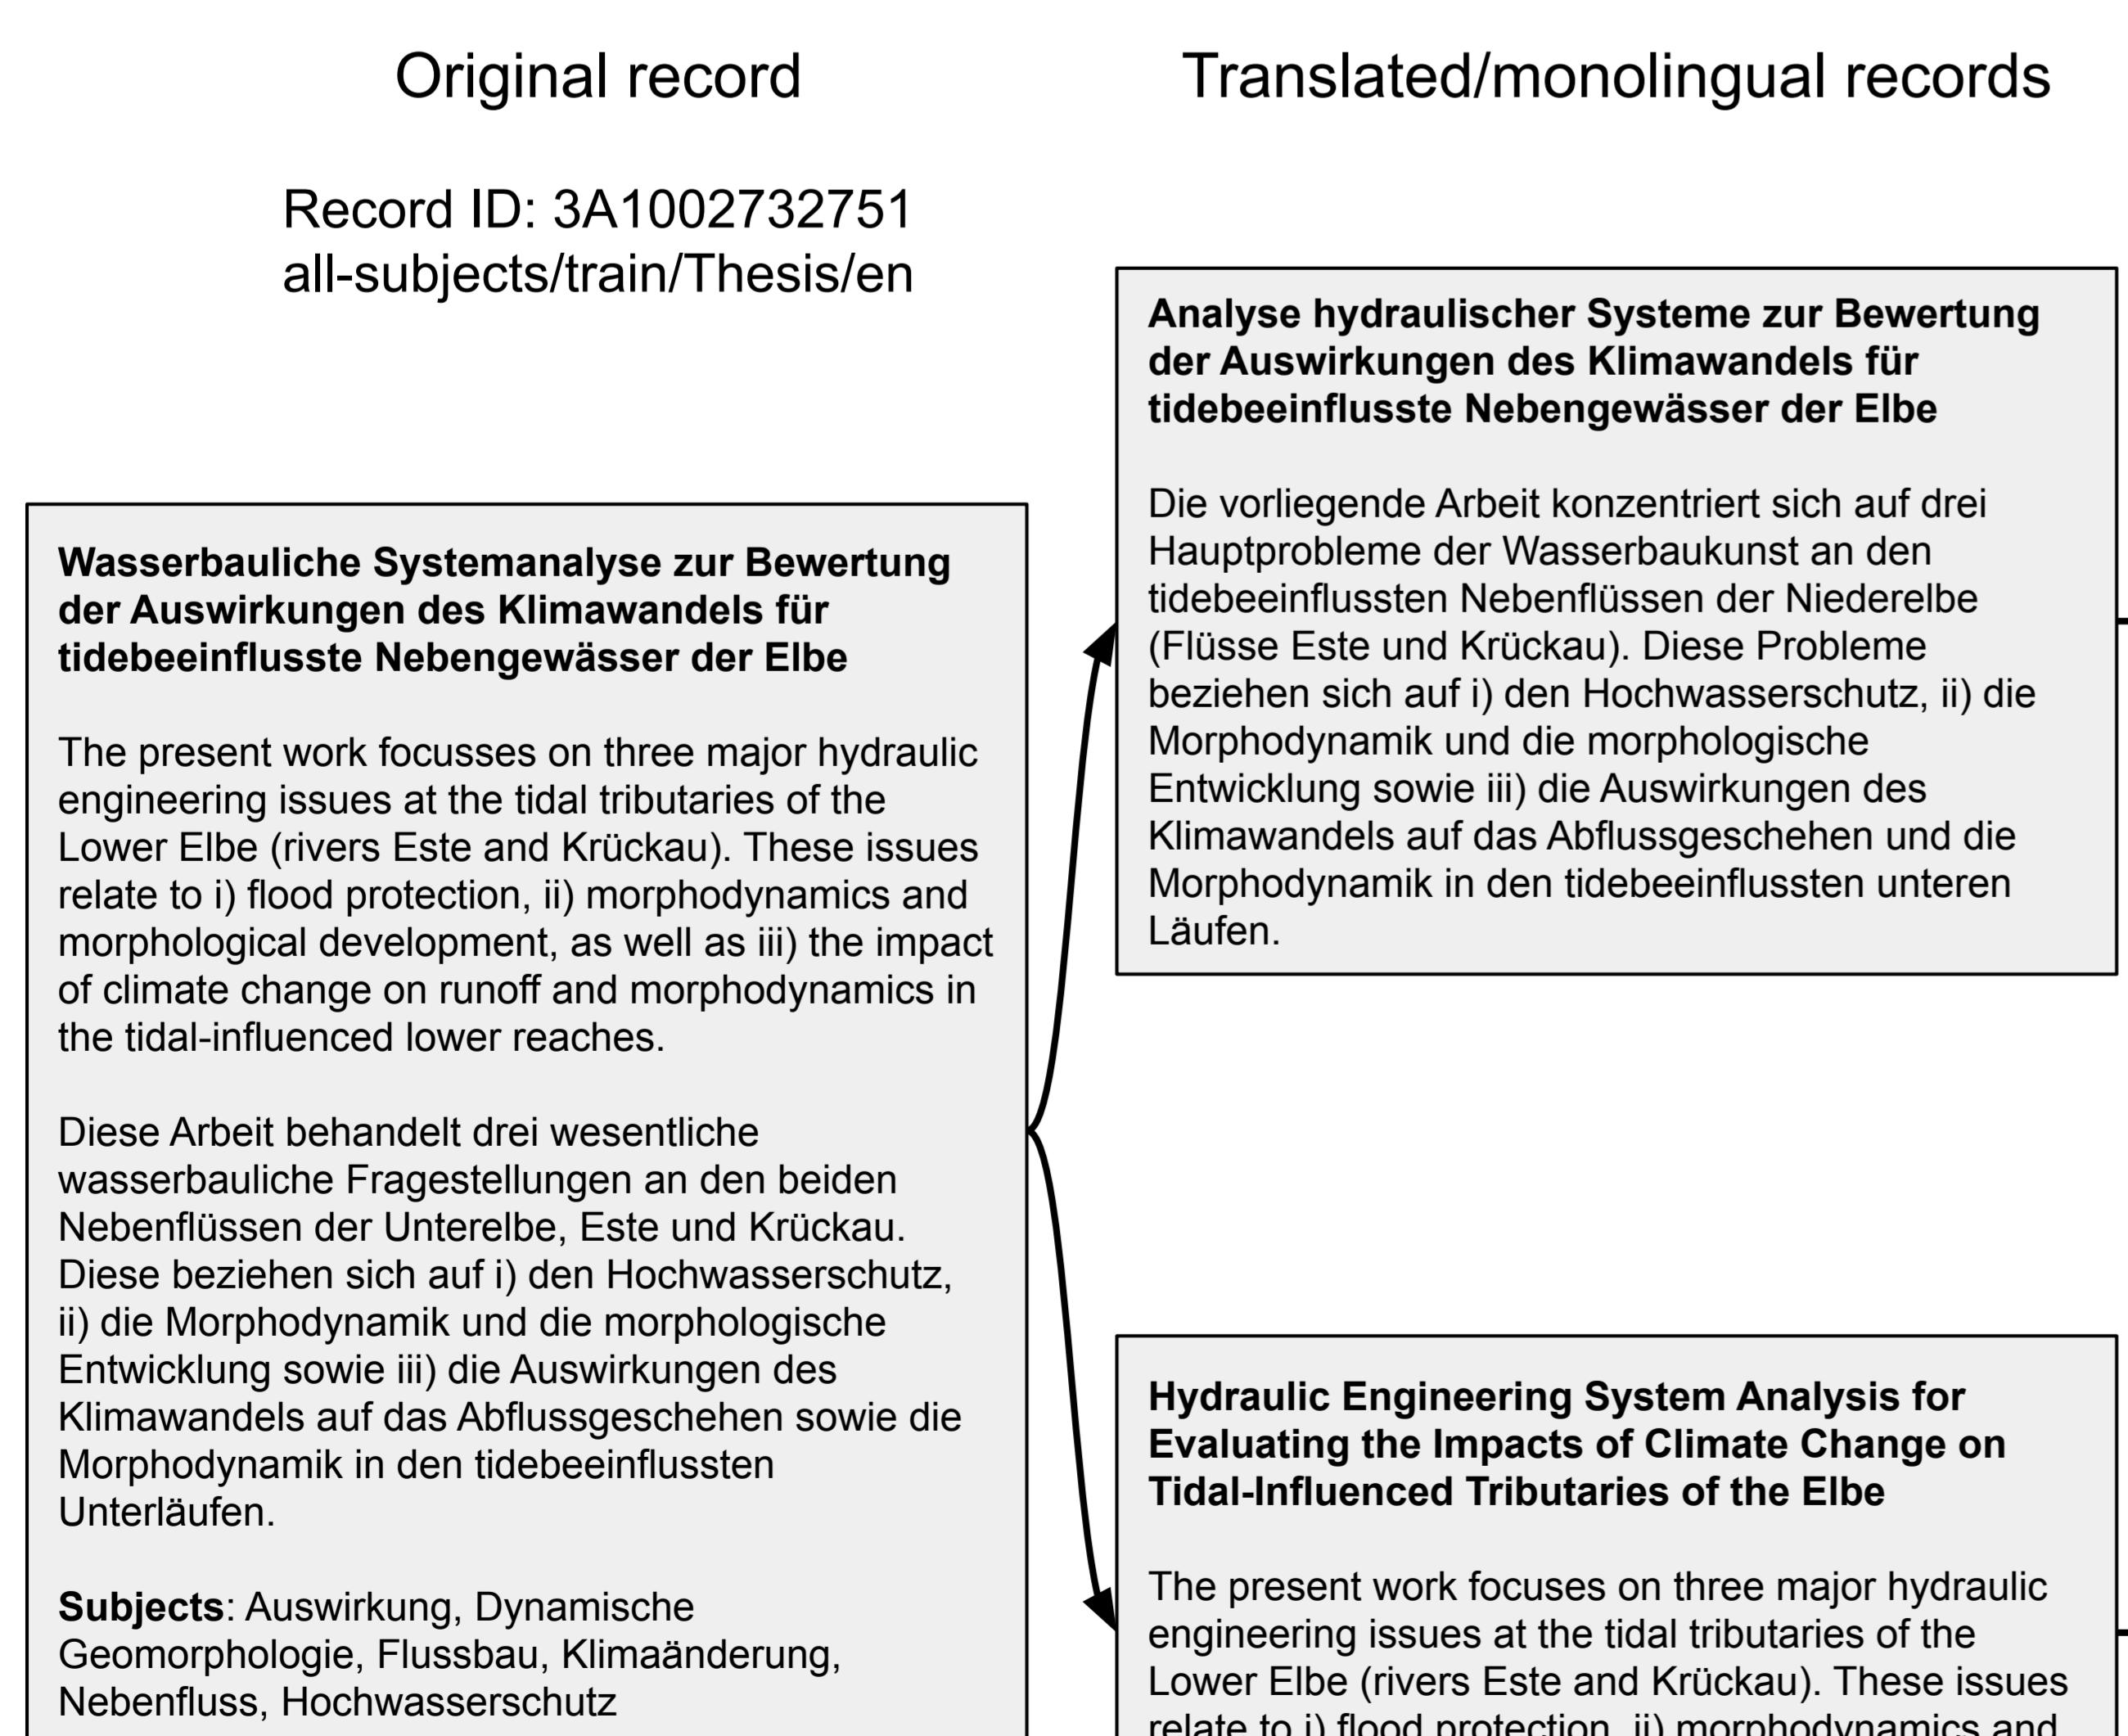


## Qualitative results

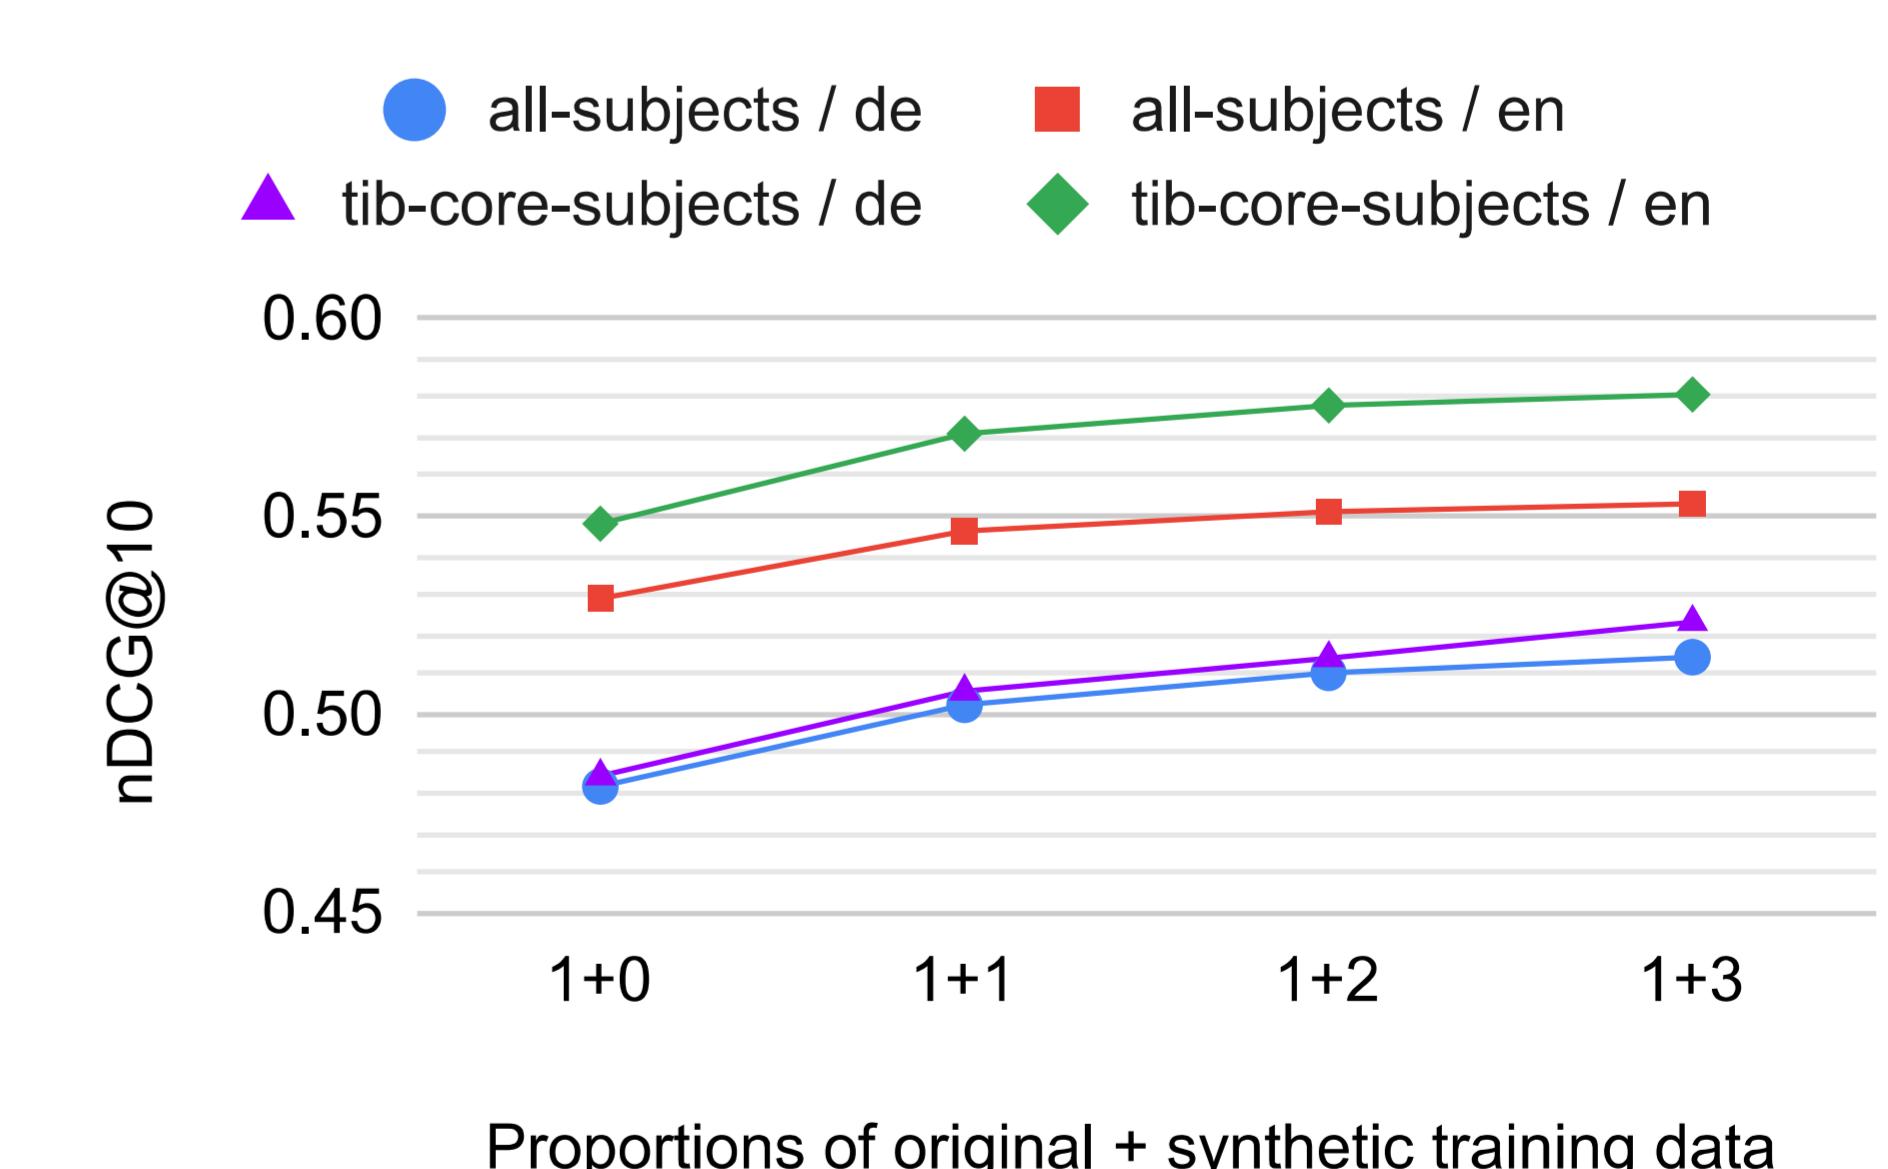
Case	System	Avg Recall	Rank
1	DNB-AI-Project	<b>0.5657</b>	1 <sup>st</sup>
	DUTIR831	0.5330	2 <sup>nd</sup>
	RUC Team	0.5199	3 <sup>rd</sup>
	<b>Annif (ours)</b>	<b>0.5024</b>	4 <sup>th</sup>
	jim	0.4928	5 <sup>th</sup>
2	DNB-AI-Project	<b>0.5094</b>	1 <sup>st</sup>
	DUTIR831	0.4851	2 <sup>nd</sup>
	RUC Team	0.4645	3 <sup>rd</sup>
	<b>Annif (ours)</b>	<b>0.4484</b>	4 <sup>th</sup>
	jim	0.4258	5 <sup>th</sup>

Table 2: Qualitative evaluation results for the top 5 teams in evaluation cases 1 and 2.

## Pre-processing example



## Synthetic training data



arxiv.org/abs/2504.19675